

```
In [28]: import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline

fcfs = pd.read_csv('Sim/simdata_fcfs.csv')
srtf = pd.read_csv('Sim/simdata_srtf.csv')
hrrn = pd.read_csv('Sim/simdata_hrrn.csv')
rr1 = pd.read_csv('Sim/simdata_rr_01.csv')
rr2 = pd.read_csv('Sim/simdata_rr_02.csv')

print(fcfs)
```

	arrival_rate	service rate	avg turnaround	total throughput	cpu util \
0	10	25.0	0.066880	10.234223	0.404404
1	11	25.0	0.071106	10.909811	0.434434
2	12	25.0	0.074153	11.762535	0.465401
3	13	25.0	0.080276	12.909523	0.516144
4	14	25.0	0.089149	13.830746	0.554534
5	15	25.0	0.099613	14.781124	0.588801
6	16	25.0	0.109013	15.976153	0.637995
7	17	25.0	0.115242	16.958142	0.674006
8	18	25.0	0.138772	17.962208	0.720504
9	19	25.0	0.151015	19.053566	0.747921
10	20	25.0	0.180845	19.712473	0.784614
11	21	25.0	0.224545	20.847494	0.837035
12	22	25.0	0.375193	21.957245	0.865872
13	23	25.0	0.490250	23.211863	0.928797
14	24	25.0	0.693975	23.699073	0.944296
15	25	25.0	3.351606	24.431284	0.988420
16	26	25.0	12.143129	25.057347	0.999355
17	27	25.0	12.808251	25.303036	0.997327
18	28	25.0	21.261573	24.960237	0.999945
19	29	25.0	22.022796	25.418182	0.999813
20	30	25.0	33.225825	25.194383	0.999811

	avg#processes
0	1.292400
1	1.345031
2	1.405956
3	1.505349
4	1.666200
5	1.880512
6	2.104179
7	2.253275
8	2.742526
9	3.122500
10	3.780988
11	4.786321
12	8.551600
13	11.361255
14	16.626306
15	83.607668
16	321.113669
17	346.004948
18	591.896172
19	634.950393
20	1001.412828

```
In [30]: x = fcfs['arrival_rate']
avg_turnaround_fcfs = fcfs['avg turnaround']
total_throughput_fcfs = fcfs['total throughput']
cpu_util_fcfs = fcfs['cpu util']
avg_processes_fcfs = fcfs['avg#processes']


avg_turnaround_srtf = srtf['avg turnaround']
total_throughput_srtf = srtf['total throughput']
cpu_util_srtf = srtf['cpu util']
avg_processes_srtf = srtf['avg#processes']

avg_turnaround_hrrn = hrrn['avg turnaround']
total_throughput_hrrn = hrrn['total throughput']
cpu_util_hrrn = hrrn['cpu util']
avg_processes_hrrn = hrrn['avg#processes']

avg_turnaround_rr1 = rr1['avg turnaround']
total_throughput_rr1 = rr1['total throughput']
cpu_util_rr1 = rr1['cpu util']
avg_processes_rr1 = rr1['avg#processes']

avg_turnaround_rr2 = rr2['avg turnaround']
total_throughput_rr2 = rr2['total throughput']
cpu_util_rr2 = rr2['cpu util']
avg_processes_rr2 = rr2['avg#processes']


plt.plot(x, avg_turnaround_fcfs,color='b')
plt.plot(x, avg_turnaround_srtf,color='g')
plt.plot(x, avg_turnaround_hrrn,color='r')
plt.plot(x, avg_turnaround_rr1,color='orange')
plt.plot(x, avg_turnaround_rr2,color='black')
location = 0 # For the best location
legend_drawn_flag = True
plt.legend(['fcfs", "srtf", "hrrn", "rr0.1", "rr.2"], loc=0, frameon=legend_drawn_flag)
plt.xlabel('arrival rate')
plt.ylabel('avg turnaround time')
plt.show()
```



AVG TURNAROUND TIME

As would be expected, SRTF has the lowest turnaround time. It's surprising, however, how much better it is than the other algorithms. HRRN also performs significantly better than the other 3.


```
In [31]: plt.plot(x, total_throughput_fcfs,color='b')
plt.plot(x, total_throughput_srtf,color='g')
plt.plot(x, total_throughput_hrrn,color='r')
plt.plot(x, total_throughput_rr1,color='orange')
plt.plot(x, total_throughput_rr2,color='black')
location = 0 # For the best location
legend_drawn_flag = True
plt.legend(['fcfs", "srtf", "hrrn", "rr0.1", "rr.2"], loc=0, frameon=legend_drawn_flag)
plt.xlabel('arrival rate')
plt.ylabel('total throughput')
plt.show()
```



TOTAL THROUGHPUT

SRTF and HRRN have the highest total throughput. This makes sense, they have the lowest average turnaround time. However the algorithms do not differ much in total throughput, only diverging at higher loads.

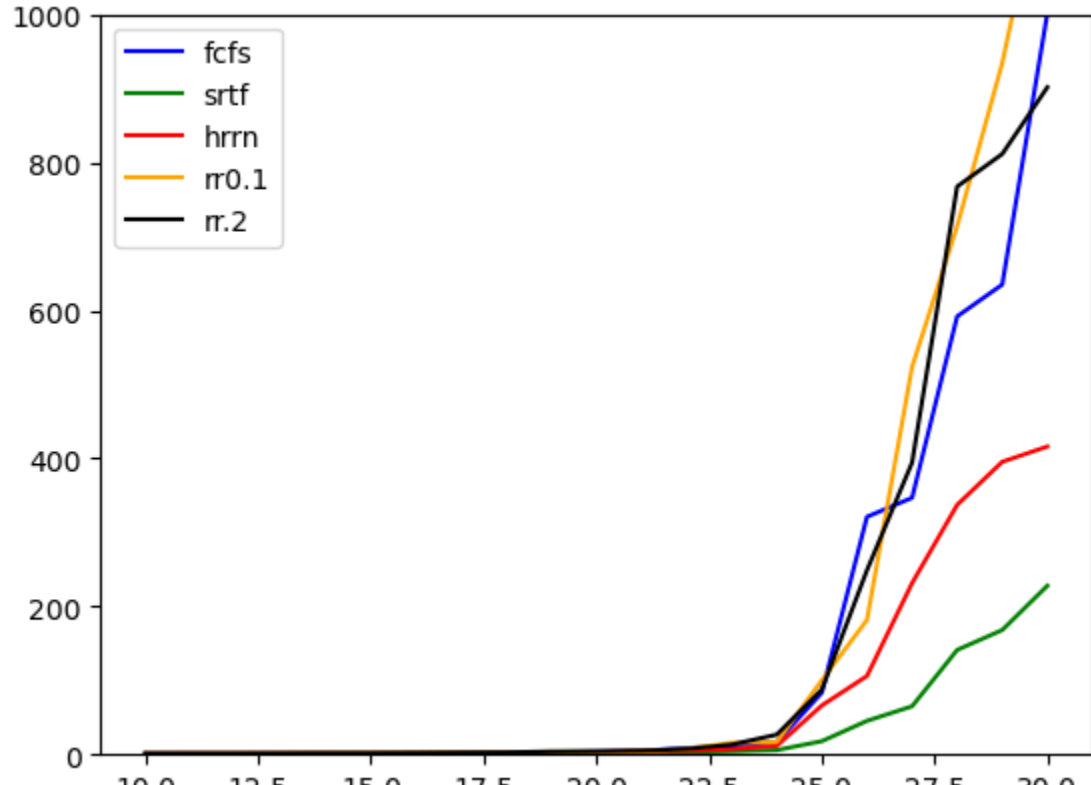
```
In [32]: plt.plot(x, cpu_util_fcfs,color='b')
plt.plot(x, cpu_util_srtf,color='g')
plt.plot(x, cpu_util_hrrn,color='r')
plt.plot(x, cpu_util_rr1,color='orange')
plt.plot(x, cpu_util_rr2,color='black')
location = 0 # For the best location
legend_drawn_flag = True
plt.legend(['fcfs", "srtf", "hrrn", "rr0.1", "rr.2"], loc=0, frameon=legend_drawn_flag)
plt.xlabel('arrival rate')
plt.ylabel('CPU Utilization')
plt.show()
```



CPU UTILIZATION

All the algorithms have similar CPU utilization. Scales linearly up to nearly 100% at high loads

```
In [21]: plt.ylim(top=1000)
plt.plot(x, avg_processes_fcfs,color='b')
plt.plot(x, avg_processes_srtf,color='g')
plt.plot(x, avg_processes_hrrn,color='r')
plt.plot(x, avg_processes_rr1,color='orange')
plt.plot(x, avg_processes_rr2,color='black')
location = 0 # For the best location
legend_drawn_flag = True
plt.legend(['fcfs", "srtf", "hrrn", "rr0.1", "rr.2"], loc=0, frameon=legend_drawn_flag)
plt.xlabel('arrival rate')
plt.ylabel('avg # of processes in ready queue')
plt.show()
```




AVG NUMBER OF PROCESSES IN READY QUEUE

SRTF does the best job at keeping the ready queue empty, followed by HRRN. The other algorithms experience skyrocketing ready queue lengths at higher loads.

Since the queue length grows quickly, lets look at the behavior at lower loads more closely.

```
In [43]: plt.plot(x, avg_processes_fcfs,color='b')
plt.plot(x, avg_processes_srtf,color='g')
plt.plot(x, avg_processes_hrrn,color='r')
plt.plot(x, avg_processes_rr1,color='orange')
plt.plot(x, avg_processes_rr2,color='black')
location = 0 # For the best location
legend_drawn_flag = True
plt.legend(['fcfs", "srtf", "hrrn", "rr0.1", "rr.2"], loc=0, frameon=legend_drawn_flag)
plt.xlabel('arrival rate')
plt.ylabel('avg # of processes in ready queue')
plt.xlim(right = 25)
plt.ylim(bottom = 0, top=100)
plt.show()
```



INTERPRETATION OF RESULTS

SRTF performs the best on all our metrics, followed by HRRN. FCFS and round robin could offer other advantages that aren't measured by this simulation.